

**Mini Project

Problem Statement: Implement a mini-project involving a real-world ML task.

Suggested topics: Weather prediction, Stock price trend analysis, Heart disease prediction, Fake news detection

15

Activities:

Data collection & cleaning

Model development

Present findings**

Heart Disease Prediction Using Machine Learning

Problem Statement

Predict whether a patient has heart disease based on clinical and demographic features such as age, cholesterol, blood pressure, and heart rate.

```
# Step 1: Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV, KFold
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Step 2: Load Dataset
# Make sure heart-disease.csv is in your current directory
df = pd.read_csv("heart-disease.csv")
```

```
# Step 3: Basic Inspection
print("✅ Dataset Loaded Successfully!")
print("Shape of data:", df.shape)
print("\nFirst 5 rows:\n", df.head())
```

✅ Dataset Loaded Successfully!
 Shape of data: (303, 14)

	age	sex	cp	trestbps	chol	fbps	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
# Step 4: Check for Missing Values
print("\nMissing Values:\n", df.isnull().sum())
```

Missing Values:

	age	sex	cp	trestbps	chol	fbps
	0	0	0	0	0	0

```
restecg    0
thalach   0
exang     0
oldpeak   0
slope     0
ca        0
thal      0
target    0
dtype: int64
```

```
# Step 5: Data Info and Statistics
print("\nData Types:\n")
print(df.info())
print("\nStatistical Summary:\n")
print(df.describe())
```


Data Types:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column   Non-Null Count Dtype  
 --- 
 0   age       303 non-null   int64  
 1   sex       303 non-null   int64  
 2   cp        303 non-null   int64  
 3   trestbps 303 non-null   int64  
 4   chol      303 non-null   int64  
 5   fbs       303 non-null   int64  
 6   restecg   303 non-null   int64  
 7   thalach   303 non-null   int64  
 8   exang     303 non-null   int64  
 9   oldpeak   303 non-null   float64 
 10  slope     303 non-null   int64  
 11  ca        303 non-null   int64  
 12  thal      303 non-null   int64  
 13  target    303 non-null   int64  
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
None
```

Statistical Summary:

	age	sex	cp	trestbps	chol	fbs	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	
	restecg	thalach	exang	oldpeak	slope	ca	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	
mean	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373	
std	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606	
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000	
50%	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000	
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000	
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000	
	thal	target					
count	303.000000	303.000000					
mean	2.313531	0.544554					
std	0.612277	0.498835					
min	0.000000	0.000000					
25%	2.000000	0.000000					
50%	2.000000	1.000000					
75%	3.000000	1.000000					
max	3.000000	1.000000					

```
# Step 6: Correlation Heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm")
plt.title("Feature Correlation Heatmap")
plt.show()
```



```
# =====
# MODEL DEVELOPMENT
# =====
# Step 7: Split Features and Target
X = df.drop('target', axis=1)
y = df['target']

# Step 8: Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Step 9: Feature Scaling
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Step 10: Initialize Models
lr = LogisticRegression(max_iter=1000, random_state=42)
rf = RandomForestClassifier(random_state=42)

# Step 11: Train Models
lr.fit(X_train, y_train)
rf.fit(X_train, y_train)
```

▼ **RandomForestClassifier** ⓘ ?
RandomForestClassifier(random_state=42)

```
# =====
# MODEL EVALUATION
# =====
```

```
# Step 12: Predictions
y_pred_lr = lr.predict(X_test)
y_pred_rf = rf.predict(X_test)

# Step 13: Evaluate Models
print("\n== Logistic Regression ==")
print("Accuracy:", round(accuracy_score(y_test, y_pred_lr), 4))
print(classification_report(y_test, y_pred_lr))

print("\n== Random Forest ==")
print("Accuracy:", round(accuracy_score(y_test, y_pred_rf), 4))
print(classification_report(y_test, y_pred_rf))

# Step 14: Confusion Matrix - Random Forest
```

```
cm = confusion_matrix(y_test, y_pred_rf)
plt.figure(figsize=(5, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix - Random Forest")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

==== Logistic Regression ====


Accuracy: 0.8033

	precision	recall	f1-score	support
0	0.86	0.68	0.76	28
1	0.77	0.91	0.83	33
accuracy				0.80
macro avg	0.82	0.79	0.80	61
weighted avg	0.81	0.80	0.80	61

==== Random Forest ====

Accuracy: 0.8361

	precision	recall	f1-score	support
0	0.95	0.68	0.79	28
1	0.78	0.97	0.86	33
accuracy				0.84
macro avg	0.87	0.82	0.83	61
weighted avg	0.86	0.84	0.83	61




```
# =====
# HYPERPARAMETER TUNING (GridSearchCV)
# =====

# Step 15: Define Parameter Grid
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}

# Step 16: Apply GridSearch with 5-Fold Cross Validation
kfold = KFold(n_splits=5, shuffle=True, random_state=42)
grid_search = GridSearchCV(
    estimator=RandomForestClassifier(random_state=42),
    param_grid=param_grid,
    cv=kfold,
    scoring='accuracy',
    n_jobs=-1,
    verbose=1
)

grid_search.fit(X_train, y_train)
```



```
# Step 17: Display Best Parameters and Accuracy
print("\nBest Parameters from Grid Search:")
print(grid_search.best_params_)
print("Best Cross-Validation Accuracy:", round(grid_search.best_score_, 4))

# Step 18: Evaluate Best Model
best_model = grid_search.best_estimator_
y_pred_best = best_model.predict(X_test)

print("\n==== Tuned Random Forest ===")
print("Accuracy:", round(accuracy_score(y_test, y_pred_best), 4))
print(classification_report(y_test, y_pred_best))
```

Best Parameters from Grid Search:
{'max_depth': 5, 'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 200}
Best Cross-Validation Accuracy: 0.8551

==== Tuned Random Forest ===
Accuracy: 0.8197

	precision	recall	f1-score	support
0	0.95	0.64	0.77	28
1	0.76	0.97	0.85	33
accuracy			0.82	61
macro avg	0.85	0.81	0.81	61
weighted avg	0.85	0.82	0.81	61

```
# =====
# FEATURE IMPORTANCE VISUALIZATION
# =====

importances = best_model.feature_importances_
features = X.columns

plt.figure(figsize=(10,6))
sns.barplot(x=importances, y=features, palette='viridis')
plt.title("Feature Importance - Tuned Random Forest")
plt.xlabel("Importance Score")
plt.ylabel("Feature")
plt.show()
```

```
/tmp/ipython-input-641666014.py:9: FutureWarning:  
  Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to  
  sns.barplot(x=importances, y=features, palette='viridis')
```

Feature Importance - Tuned Random Forest

```
# =====  
# FINAL COMPARISON  
# =====  
print("\n--- MODEL COMPARISON ---")  
print(f"Logistic Regression Accuracy: {round(accuracy_score(y_test, y_pred_lr), 4)}")  
print(f"Random Forest Accuracy: {round(accuracy_score(y_test, y_pred_rf), 4)}")  
print(f"Tuned Random Forest Accuracy: {round(accuracy_score(y_test, y_pred_best), 4)}")
```

